

# CS 5644: Assignment 4 report

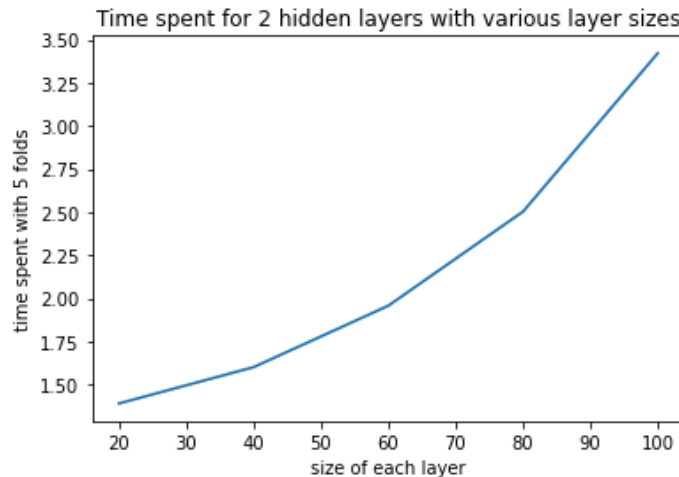
## 1.1 Question 1 Multilayer Perceptron Classifier

MLP classifier is created with the argument `max_iter = 1000` and `solver = lbfgs`. It uses `MinMaxScaler` for scaling and 5-fold for cross-validation. The following are the results of different MLP architectures in different scenarios.


```
2 hidden layers with 20 nodes in each layer (tanh):
  average precision: 0.9058907226786792
  average recall: 0.9050763852828249
  average f1: 0.904931213848568
2 hidden layers with 20 nodes in each layer (relu):
  average precision: 0.9152678185454448
  average recall: 0.9150510588283479
  average f1: 0.9147416642698127
2 hidden layers with 100 nodes in each layer (tanh):
  average precision: 0.9339836001083535
  average recall: 0.9342328743197135
  average f1: 0.9338480617257192
2 hidden layers with 100 nodes in each layer (relu):
  average precision: 0.9334280362354008
  average recall: 0.9335172902002926
  average f1: 0.9331148423479796
5 hidden layers with 20 nodes in each layer (tanh):
  average precision: 0.9091827932250277
  average recall: 0.9095224502309357
  average f1: 0.9089670266698503
5 hidden layers with 20 nodes in each layer (relu):
  average precision: 0.8990806781933156
  average recall: 0.8994058645244669
  average f1: 0.8987233476951009
5 hidden layers with 100 nodes in each layer (tanh):
  average precision: 0.9333547933234442
  average recall: 0.9334114745205362
  average f1: 0.9331093947924789
5 hidden layers with 100 nodes in each layer (relu):
  average precision: 0.9409263287085233
  average recall: 0.9411617230717602
  average f1: 0.9407422354442602
```

| Scenario \ Activation funs                   | tanh  | ReLU  |
|--|---|---|
| 2 hidden layers with 20 nodes in each layer  | precision = 0.906<br>recall = 0.905<br>f1 = 0.905 | precision = 0.915<br>recall = 0.915<br>f1 = 0.91  |
| 2 hidden layers with 100 nodes in each layer | precision = 0.934<br>recall = 0.934<br>f1 = 0.934 | precision = 0.933<br>recall = 0.934<br>f1 = 0.933 |
| 5 hidden layers with 20 nodes in each layer  | precision = 0.91<br>recall = 0.91<br>f1 = 0.91    | precision = 0.899<br>recall = 0.899<br>f1 = 0.899 |
| 5 hidden layers with 100 nodes in each layer | precision = 0.933<br>recall = 0.933<br>f1 = 0.933 | precision = 0.941<br>recall = 0.941<br>f1 = 0.941 |

Observe the time spent to fit each model as the size of layers increases for various layer sizes. 5-fold cross-validation and ReLU activation function are applied. The number of layers is set to 2. The graph of the result is presented as follows.



Observe the time spent to fit each model as the number of layers increases for various numbers of layers. 5-fold cross-validation and ReLU activation function are applied. The size of layers is set to 20. See the graph of the result below.



From the results, we see that either we increase the size of each layer or increase the number of layers, the time it takes to fit the model mounting. The reason is that the more nodes we add to the model, the more calculations are needed for training.

If we stop using MinMaxScaler for simple/smallish networks, the input variables may have different scales. In that case, large input values can result in a model that learns large weight values. Consequently, the model may suffer from poor performance during learning and sensitivity to input values resulting in higher generalization errors.

## 1.2 Question 2

When the learning rate is set to 0.5, the performance will oscillate more and converge more quickly over training epochs than when the learning rate is set to 0.001. This is because the learning rate determines how much an updating step influences the current value of the weights, that is, a larger learning rate causes a bigger divergence of weights.